The Runner class in the OpenAI Agents SDK plays a central role in making agents function. It's like the operator or engine behind the scenes that takes a user's request and tells the agent how to respond. The agent itself holds information about what it can do and how it should behave, but it's the Runner that actually runs that agent when a user interacts with the system. The main job of the Runner is to take a user prompt, combine it with the agent's instructions and tools, and then generate a response. It makes sure the agent follows its assigned role and uses the right tools to answer the user's question or perform a task. The Runner is also responsible for managing the flow of communication between the user and the agent. It passes the prompt to the agent, waits for a response, and returns that response to the user. This keeps the process smooth and controlled. In systems where multiple agents are involved, the Runner plays an even bigger role by handling handoffs between agents. For example, if a general assistant agent receives a prompt but realizes it's a billing issue, the Runner can pass the task to a billing-specific agent who is better equipped to handle it. This way, each agent focuses on what it's best at, and the Runner ensures the right agent is in charge at the right time.

Another important purpose of the Runner is to keep the agent's definition separate from its execution. The Agent class defines what the agent is—its instructions, goals, and tools—while the Runner takes care of actually using that agent in real situations. This separation makes it easy to reuse the same agent in different workflows or scenarios without rewriting its definition. If you want to run the same agent on different types of prompts or in different applications, you just use the Runner with different inputs. The Runner also manages the entire process during execution. It starts the task, makes sure the agent processes everything correctly, and handles any errors that occur. For example, if the agent can't understand the prompt or if something goes wrong, the Runner knows how to catch those errors and either handle them or report them properly. This helps the system stay stable even when there are problems or unexpected inputs.

In multi-agent systems, where agents work together like a team, the Runner coordinates how these agents interact. It ensures that each task is passed to the correct agent and that agents can collaborate smoothly. This is helpful in more advanced use cases like customer service, where one user query might be handled by multiple agents—such as a general support agent, a billing agent, and a tech support agent. The Runner makes sure each agent gets involved at the right time without confusion. The design of the Runner class also reflects the SDK's lightweight and user-friendly approach. It's made to be simple for developers to use. For instance, developers can just call a method like `Runner.run(agent, prompt)` and the Runner handles the entire workflow behind the scenes. This makes it easier to build real applications like chatbots, virtual assistants, and automated customer service tools without needing deep expertise in managing AI systems.

In summary, the Runner class is crucial in the OpenAI Agents SDK because it brings agents to life. It runs them, manages communication, handles teamwork between agents, and keeps everything organized and error-free. Without the Runner, agents would just be definitions without action. Thanks to the Runner, agents can respond to user prompts, work together on tasks, and function in real-world applications efficiently and reliably.